

How to Read 100 Million Blogs (& Classify Deaths Without Physicians)

Gary King
Harvard University

April 11, 2007

- Daniel Hopkins and Gary King. “Extracting Systematic Social Science Meaning from Text”

- Daniel Hopkins and Gary King. “Extracting Systematic Social Science Meaning from Text”
- Gary King and Ying Lu. “Verbal Autopsy Methods with Multiple Causes of Death,” tentatively to appear, *Statistical Science*

- Daniel Hopkins and Gary King. “Extracting Systematic Social Science Meaning from Text”
- Gary King and Ying Lu. “Verbal Autopsy Methods with Multiple Causes of Death,” tentatively to appear, *Statistical Science*
- Copies at <http://gking.harvard.edu>

Content Analysis: Past and Future

Content Analysis: Past and Future

- Dates to the 1600s: The Church tracked nonreligious texts by classifying newspaper stories

Content Analysis: Past and Future

- Dates to the 1600s: The Church tracked nonreligious texts by classifying newspaper stories
- Prominent early social scientists used it: Berelson, de Grazia, etc.

Content Analysis: Past and Future

- Dates to the 1600s: The Church tracked nonreligious texts by classifying newspaper stories
- Prominent early social scientists used it: Berelson, de Grazia, etc.
- Spread to vast array of fields (use increased six-fold 1980–2000)

Content Analysis: Past and Future

- Dates to the 1600s: The Church tracked nonreligious texts by classifying newspaper stories
- Prominent early social scientists used it: Berelson, de Grazia, etc.
- Spread to vast array of fields (use increased six-fold 1980–2000)
- New applications: explosive increase in web pages, blogs, emails, digitized books and articles, audio recordings (automatically converted to text), and government reports, legislative hearings and records, electronic medical records, etc.

Content Analysis: Past and Future

- Dates to the 1600s: The Church tracked nonreligious texts by classifying newspaper stories
- Prominent early social scientists used it: Berelson, de Grazia, etc.
- Spread to vast array of fields (use increased six-fold 1980–2000)
- New applications: explosive increase in web pages, blogs, emails, digitized books and articles, audio recordings (automatically converted to text), and government reports, legislative hearings and records, electronic medical records, etc.
- Infeasible to expand hand coding efforts much further

Content Analysis: Past and Future

- Dates to the 1600s: The Church tracked nonreligious texts by classifying newspaper stories
- Prominent early social scientists used it: Berelson, de Grazia, etc.
- Spread to vast array of fields (use increased six-fold 1980–2000)
- New applications: explosive increase in web pages, blogs, emails, digitized books and articles, audio recordings (automatically converted to text), and government reports, legislative hearings and records, electronic medical records, etc.
- Infeasible to expand hand coding efforts much further
- Automated methods are essential

Inputs and Target Quantities of Interest

Inputs and Target Quantities of Interest

- Available inputs:

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest
 - individual document classification

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest
 - individual document classification
 - proportion of documents in each category

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest
 - individual document classification
 - proportion of documents in each category
 - *Can* get the 2nd by aggregating the 1st (turns out not to be necessary!)

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest
 - individual document classification
 - proportion of documents in each category
 - *Can* get the 2nd by aggregating the 1st (turns out not to be necessary!)
 - E.g., classify constituents' letters to a member of congress by policy area, or estimate proportion of letters in each policy area

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest
 - individual document classification
 - proportion of documents in each category
 - *Can* get the 2nd by aggregating the 1st (turns out not to be necessary!)
 - E.g., classify constituents' letters to a member of congress by policy area, or estimate proportion of letters in each policy area
 - E.g., classify emails as spam or not, or estimate proportion of email that is spam

Inputs and Target Quantities of Interest

- Available inputs:
 - Large set of text documents
 - A set of (mutually exclusive and exhaustive) categories
 - A small subset of documents hand-coded into the categories
- Quantities of interest
 - individual document classification
 - proportion of documents in each category
 - *Can* get the 2nd by aggregating the 1st (turns out not to be necessary!)
 - E.g., classify constituents' letters to a member of congress by policy area, or estimate proportion of letters in each policy area
 - E.g., classify emails as spam or not, or estimate proportion of email that is spam
- Maximizing one goal won't get you the other: high classification accuracy can coexist with huge biases in category proportions

Our Approach

Our Approach

- Gives unbiased estimates of population proportions

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method
- No problem if classification accuracy is low

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method
- No problem if classification accuracy is low
- (And individual classification is not necessary)

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method
- No problem if classification accuracy is low
- (And individual classification is not necessary)
- No parametric modeling assumptions

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method
- No problem if classification accuracy is low
- (And individual classification is not necessary)
- No parametric modeling assumptions
- The hand coded subset need not be a random sample

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method
- No problem if classification accuracy is low
- (And individual classification is not necessary)
- No parametric modeling assumptions
- The hand coded subset need not be a random sample
- Scales to large numbers of documents

Our Approach

- Gives unbiased estimates of population proportions
- Works better than aggregating the best classification method
- No problem if classification accuracy is low
- (And individual classification is not necessary)
- No parametric modeling assumptions
- The hand coded subset need not be a random sample
- Scales to large numbers of documents
- Separately: propose correction for imperfect inter-coder reliability (i.e., should work better than hand coding everything if that were feasible)

Blogs as a Running Example

Blogs as a Running Example

- Blogs (web logs): web version of a daily diary, with posts listed in reverse chronological order.

Blogs as a Running Example

- Blogs (web logs): web version of a daily diary, with posts listed in reverse chronological order.
- 8% of U.S. Internet users (12 million) have blogs

Blogs as a Running Example

- Blogs (web logs): web version of a daily diary, with posts listed in reverse chronological order.
- 8% of U.S. Internet users (12 million) have blogs
- Growth: ≈ 0 in 2000 to 39–100 million worldwide now.

Blogs as a Running Example

- Blogs (web logs): web version of a daily diary, with posts listed in reverse chronological order.
- 8% of U.S. Internet users (12 million) have blogs
- Growth: ≈ 0 in 2000 to 39–100 million worldwide now.
- A democratic technology: 6 million in China and 700,000 in Iran(!)

Blogs as a Running Example

- Blogs (web logs): web version of a daily diary, with posts listed in reverse chronological order.
- 8% of U.S. Internet users (12 million) have blogs
- Growth: ≈ 0 in 2000 to 39–100 million worldwide now.
- A democratic technology: 6 million in China and 700,000 in Iran(!)
- “We are living through the largest expansion of expressive capability in the history of the human race”

One specific quantity of interest

One specific quantity of interest

- Subject: the grand conversation about the American presidency

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates

- Specific categories:

<u>Label</u>	<u>Category</u>
-2	extremely negative
-1	negative
0	neutral
1	positive
2	extremely positive
NA	no opinion expressed
NB	not a blog

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates
- Specific categories:

<u>Label</u>	<u>Category</u>
-2	extremely negative
-1	negative
0	neutral
1	positive
2	extremely positive
NA	no opinion expressed
NB	not a blog
- Hard case:

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates
- Specific categories:

<u>Label</u>	<u>Category</u>
-2	extremely negative
-1	negative
0	neutral
1	positive
2	extremely positive
NA	no opinion expressed
NB	not a blog
- Hard case:
 - Part ordinal, part nominal categorization

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates
- Specific categories:

<u>Label</u>	<u>Category</u>
-2	extremely negative
-1	negative
0	neutral
1	positive
2	extremely positive
NA	no opinion expressed
NB	not a blog
- Hard case:
 - Part ordinal, part nominal categorization
 - “Sentiment categorization is more difficult than topic classification”

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates
- Specific categories:

<u>Label</u>	<u>Category</u>
-2	extremely negative
-1	negative
0	neutral
1	positive
2	extremely positive
NA	no opinion expressed
NB	not a blog
- Hard case:
 - Part ordinal, part nominal categorization
 - “Sentiment categorization is more difficult than topic classification”
 - Language ranges from “my crunchy gf thinks dubya hid the wmd’s, :)!” to the Queen’s English

One specific quantity of interest

- Subject: the grand conversation about the American presidency
- Question: affect about President Bush and 2008 candidates

- Specific categories:

<u>Label</u>	<u>Category</u>
-2	extremely negative
-1	negative
0	neutral
1	positive
2	extremely positive
NA	no opinion expressed
NB	not a blog

- Hard case:
 - Part ordinal, part nominal categorization
 - “Sentiment categorization is more difficult than topic classification”
 - Language ranges from “my crunchy gf thinks dubya hid the wmd’s, :)!” to the Queen’s English
 - Little common internal structure (no inverted pyramid)

The Conversation about John Kerry's Botched Joke


The Conversation about John Kerry's Botched Joke

You know, education — if you make the most of it . . . you can do well. If you don't, you get stuck in Iraq.

The Conversation about John Kerry's Botched Joke

You know, education — if you make the most of it . . . you can do well. If you don't, you get stuck in Iraq.

Affect Towards John Kerry



2006–2007

Representing Text as Numbers

Representing Text as Numbers

- **Filter:** choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.

Representing Text as Numbers

- **Filter:** choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.
- **Preprocess:** convert to lower case, remove punctuation, perform stemming (reduce “consist”, “consisted”, “consistency”, “consistent”, “consistently”, “consisting”, and “consists”, to their stem: “consist”)

Representing Text as Numbers

- **Filter**: choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.
- **Preprocess**: convert to lower case, remove punctuation, perform stemming (reduce “consist”, “consisted”, “consistency”, “consistent”, “consistently”, “consisting”, and “consists”, to their stem: “consist”)
- **Code variables** as presence or absence of unique unigrams, bigrams, trigrams, etc.

Representing Text as Numbers

- **Filter**: choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.
- **Preprocess**: convert to lower case, remove punctuation, perform stemming (reduce “consist”, “consisted”, “consistency”, “consistent”, “consistently”, “consisting”, and “consists”, to their stem: “consist”)
- **Code variables** as presence or absence of unique unigrams, bigrams, trigrams, etc.
- **Example**:

Representing Text as Numbers

- **Filter:** choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.
- **Preprocess:** convert to lower case, remove punctuation, perform stemming (reduce “consist”, “consisted”, “consistency”, “consistent”, “consistently”, “consisting”, and “consists”, to their stem: “consist”)
- **Code variables** as presence or absence of unique unigrams, bigrams, trigrams, etc.
- **Example:**
 - Our 10,771 blog posts about Bush and Clinton:
201,676 unigrams, 2,392,027 bigrams, 5,761,979 trigrams.

Representing Text as Numbers

- **Filter**: choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.
- **Preprocess**: convert to lower case, remove punctuation, perform stemming (reduce “consist”, “consisted”, “consistency”, “consistent”, “consistently”, “consisting”, and “consists”, to their stem: “consist”)
- **Code variables** as presence or absence of unique unigrams, bigrams, trigrams, etc.
- **Example**:
 - Our 10,771 blog posts about Bush and Clinton:
201,676 unigrams, 2,392,027 bigrams, 5,761,979 trigrams.
 - Unigrams in $> 1\%$ or $< 99\%$ of documents: 3,672 variables

Representing Text as Numbers

- **Filter:** choose English language blogs that mention Bush (“Bush”, “George W.”, “Dubya”, “King George”, etc.), Hillary Clinton (“Senator Clinton”, “Hillary”, “Hitlery”, “Mrs. Clinton”), etc.
- **Preprocess:** convert to lower case, remove punctuation, perform stemming (reduce “consist”, “consisted”, “consistency”, “consistent”, “consistently”, “consisting”, and “consists”, to their stem: “consist”)
- **Code variables** as presence or absence of unique unigrams, bigrams, trigrams, etc.
- **Example:**
 - Our 10,771 blog posts about Bush and Clinton:
201,676 unigrams, 2,392,027 bigrams, 5,761,979 trigrams.
 - Unigrams in $> 1\%$ or $< 99\%$ of documents: 3,672 variables
 - Groups infinite possible posts into “only” $2^{3,672}$ distinct types

Notation

- Document Category

$$D_i = \begin{cases} -2 & \text{extremely negative} \\ -1 & \text{negative} \\ 0 & \text{neutral} \\ 1 & \text{positive} \\ 2 & \text{extremely positive} \\ \text{NA} & \text{no opinion expressed} \\ \text{NB} & \text{not a blog} \end{cases}$$

- Document Category

$$D_i = \begin{cases} -2 & \text{extremely negative} \\ -1 & \text{negative} \\ 0 & \text{neutral} \\ 1 & \text{positive} \\ 2 & \text{extremely positive} \\ \text{NA} & \text{no opinion expressed} \\ \text{NB} & \text{not a blog} \end{cases}$$

- Word Stem Profile:

$$S_i = \begin{cases} S_{i1} = 1 & \text{if "awful" is used, 0 if not} \\ S_{i2} = 1 & \text{if "good" is used, 0 if not} \\ \vdots & \vdots \\ S_{iK} = 1 & \text{if "except" is used, 0 if not} \end{cases}$$

Quantities of Interest

Quantities of Interest

- Computer Science: individual document **classifications**

$$D_1, D_2, \dots, D_L$$

Quantities of Interest

- Computer Science: individual document **classifications**

$$D_1, D_2, \dots, D_L$$

- Social Science: **proportions** in each category

$$P(D) = \begin{pmatrix} P(D = -2) \\ P(D = -1) \\ P(D = 0) \\ P(D = 1) \\ P(D = 2) \\ P(D = \text{NA}) \\ P(D = \text{NB}) \end{pmatrix}$$

Issues with Existing Statistical Approaches

Issues with Existing Statistical Approaches

1 Direct Sampling

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

② Aggregation of model-based individual classifications

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

② Aggregation of model-based individual classifications

- Biased if not random sample

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

② Aggregation of model-based individual classifications

- Biased if not random sample
- Models $P(D|\mathbf{S})$, but the world works as $P(\mathbf{S}|D)$

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

② Aggregation of model-based individual classifications

- Biased if not random sample
- Models $P(D|\mathbf{S})$, but the world works as $P(\mathbf{S}|D)$
- Bias unless

Issues with Existing Statistical Approaches

1 Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

2 Aggregation of model-based individual classifications

- Biased if not random sample
- Models $P(D|\mathbf{S})$, but the world works as $P(\mathbf{S}|D)$
- Bias unless
 - $P(D|\mathbf{S})$ encompasses the “true” model.

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

② Aggregation of model-based individual classifications

- Biased if not random sample
- Models $P(D|\mathbf{S})$, but the world works as $P(\mathbf{S}|D)$
- Bias unless
 - $P(D|\mathbf{S})$ encompasses the “true” model.
 - $\mathbf{S}$ spans the space of all predictors of D (i.e., all information in the document)

Issues with Existing Statistical Approaches

① Direct Sampling

- Classification of population documents not necessary
- Biased without a random sample
- nonrandomness common due to population drift, studying data subdivisions, etc.

② Aggregation of model-based individual classifications

- Biased if not random sample
- Models $P(D|\mathbf{S})$, but the world works as $P(\mathbf{S}|D)$
- Bias unless
 - $P(D|\mathbf{S})$ encompasses the “true” model.
 - $\mathbf{S}$ spans the space of all predictors of D (i.e., all information in the document)
- Bias even with optimal classification and high % correctly classified

Using Misclassification Rates to Correct Proportions

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**
- Use labeled set to **estimate misclassification rates** (by cross-validation)

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**
- Use labeled set to **estimate misclassification rates** (by cross-validation)
- **Aggregate classifications** to category proportions

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**
- Use labeled set to **estimate misclassification rates** (by cross-validation)
- **Aggregate classifications** to category proportions
- **Use misclassification rates to correct proportions**

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**
- Use labeled set to **estimate misclassification rates** (by cross-validation)
- **Aggregate classifications** to category proportions
- **Use misclassification rates to correct proportions**
- **Result:** vastly improved estimates of category proportions

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**
- Use labeled set to **estimate misclassification rates** (by cross-validation)
- **Aggregate classifications** to category proportions
- **Use misclassification rates to correct proportions**
- **Result:** vastly improved estimates of category proportions
- (Assumes misclassification rates are estimated well)

Using Misclassification Rates to Correct Proportions

- Use some method to **classify unlabeled documents**
- Use labeled set to **estimate misclassification rates** (by cross-validation)
- **Aggregate classifications** to category proportions
- **Use misclassification rates to correct proportions**
- **Result:** vastly improved estimates of category proportions
- (Assumes misclassification rates are estimated well)
- (still requires individual classification)

Formalization from Epidemiology

(Levy and Kass, 1970)

Formalization from Epidemiology

(Levy and Kass, 1970)

- Accounting identity for 2 categories:

$$P(\hat{D} = 1) = (\text{sens})P(D = 1) + (1 - \text{spec})P(D = 2)$$

Formalization from Epidemiology

(Levy and Kass, 1970)

- Accounting identity for 2 categories:

$$P(\hat{D} = 1) = (\text{sens})P(D = 1) + (1 - \text{spec})P(D = 2)$$

- Solve:

$$P(D = 1) = \frac{P(\hat{D} = 1) - (1 - \text{spec})}{\text{sens} - (1 - \text{spec})}$$

Formalization from Epidemiology

(Levy and Kass, 1970)

- Accounting identity for 2 categories:

$$P(\hat{D} = 1) = (\text{sens})P(D = 1) + (1 - \text{spec})P(D = 2)$$

- Solve:

$$P(D = 1) = \frac{P(\hat{D} = 1) - (1 - \text{spec})}{\text{sens} - (1 - \text{spec})}$$

- Use this equation to correct $P(\hat{D})$

Generalizations: J Categories, No Individual Classification

(King and Lu, 2007)

Generalizations: J Categories, No Individual Classification

(King and Lu, 2007)

- Accounting identity for J categories

$$P(\hat{D} = j) = \sum_{j'=1}^J P(\hat{D} = j | D = j') P(D = j')$$

Generalizations: J Categories, No Individual Classification

(King and Lu, 2007)

- Accounting identity for J categories

$$P(\hat{D} = j) = \sum_{j'=1}^J P(\hat{D} = j | D = j') P(D = j')$$

- Drop $\hat{D}$ calculation, since $\hat{D} = f(\mathbf{S})$:

$$P(\mathbf{S} = s) = \sum_{j'=1}^J P(\mathbf{S} = s | D = j') P(D = j')$$

Generalizations: J Categories, No Individual Classification

(King and Lu, 2007)

- Accounting identity for J categories

$$P(\hat{D} = j) = \sum_{j'=1}^J P(\hat{D} = j | D = j') P(D = j')$$

- Drop $\hat{D}$ calculation, since $\hat{D} = f(\mathbf{S})$:

$$P(\mathbf{S} = s) = \sum_{j'=1}^J P(\mathbf{S} = s | D = j') P(D = j')$$

- Simplify to an equivalent matrix expression:

$$P(\mathbf{S}) = P(\mathbf{S} | D) P(D)$$

Estimation

The matrix expression again:

$$\underset{2^K \times 1}{P(\mathbf{S})} = \underset{2^K \times J}{P(\mathbf{S}|D)} \underset{J \times 1}{P(D)}$$

Estimation

The matrix expression again:

$$\underset{2^K \times 1}{P(\mathbf{S})} = \underset{2^K \times J}{P(\mathbf{S}|D)} \underset{J \times 1}{P(D)}$$

Document category proportions (quantity of interest)

Estimation

The matrix expression again:

$$\underset{2^K \times 1}{P(\mathbf{S})} = \underset{2^K \times J}{P(\mathbf{S}|D)} \underset{J \times 1}{P(D)}$$

Word stem profile proportions (estimate in unlabeled set by tabulation)

Estimation

The matrix expression again:

$$\underset{2^K \times 1}{P(\mathbf{S})} = \underset{2^K \times J}{P(\mathbf{S}|D)} \underset{J \times 1}{P(D)}$$

Word stem profiles, by category (estimate in *labeled* set by tabulation)

Estimation

The matrix expression again:

$$\begin{matrix} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J \quad J \times 1 \end{matrix}$$
$$\implies \mathbf{Y} = \mathbf{X}\beta$$

Alternative symbols (to emphasize the linear equation)

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 & \\ \implies \mathbf{Y} = \mathbf{X}\beta & \implies & \beta = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \end{array}$$

Solve for quantity of interest (with no error term)

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D) & P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 & \\ \implies \mathbf{Y} = \mathbf{X}\beta & \implies & \beta = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} & & \end{array}$$

- Technical estimation issues:

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D) & P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 & \\ \implies \mathbf{Y} = \mathbf{X}\beta & \implies & \beta = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \end{array}$$

- Technical estimation issues:
 - 2^K is enormous, far larger than any existing computer

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D) & P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 \\ \implies \mathbf{Y} = \mathbf{X}\beta & \implies & \beta = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \end{array}$$

- Technical estimation issues:
 - 2^K is enormous, far larger than any existing computer
 - $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 \\ \implies \mathbf{Y} = \mathbf{X}\beta & \implies & \beta = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \end{array}$$

- Technical estimation issues:
 - 2^K is enormous, far larger than any existing computer
 - $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse
 - Elements of $P(D)$ must be between 0 and 1 and sum to 1

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D) & P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 & \\ \implies \mathbf{Y} = \mathbf{X}\beta & \implies & \beta = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} & & \end{array}$$

- Technical estimation issues:
 - 2^K is enormous, far larger than any existing computer
 - $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse
 - Elements of $P(D)$ must be between 0 and 1 and sum to 1
- Solutions

Estimation

The matrix expression again:

$$\begin{array}{ccc} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J \quad J \times 1 \end{array}$$
$$\implies Y = X\beta \quad \implies \quad \beta = (X'X)^{-1}X'y$$

- Technical estimation issues:
 - 2^K is enormous, far larger than any existing computer
 - $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse
 - Elements of $P(D)$ must be between 0 and 1 and sum to 1
- Solutions
 - Use subsets of $\mathbf{S}$; average results

Estimation

The matrix expression again:

$$\begin{array}{ccccc} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J & J \times 1 \\ \implies Y = X\beta & \implies & \beta = (X'X)^{-1}X'y \end{array}$$

- Technical estimation issues:
 - 2^K is enormous, far larger than any existing computer
 - $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse
 - Elements of $P(D)$ must be between 0 and 1 and sum to 1
- Solutions
 - Use subsets of $\mathbf{S}$; average results
 - Equivalent to kernel density smoothing of sparse categorical data

Estimation

The matrix expression again:

$$\begin{array}{ccc} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J \quad J \times 1 \end{array}$$
$$\implies Y = X\beta \implies \beta = (X'X)^{-1}X'y$$

- Technical estimation issues:

- 2^K is enormous, far larger than any existing computer
- $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse
- Elements of $P(D)$ must be between 0 and 1 and sum to 1

- Solutions

- Use subsets of $\mathbf{S}$; average results
- Equivalent to kernel density smoothing of sparse categorical data
- Use constrained LS to constrain $P(D)$ to simplex

Estimation

The matrix expression again:

$$\begin{array}{ccc} P(\mathbf{S}) & = & P(\mathbf{S}|D)P(D) \\ 2^K \times 1 & & 2^K \times J \quad J \times 1 \end{array}$$
$$\implies Y = X\beta \quad \implies \quad \beta = (X'X)^{-1}X'y$$

- Technical estimation issues:

- 2^K is enormous, far larger than any existing computer
- $P(\mathbf{S})$ and $P(\mathbf{S}|D)$ will be too sparse
- Elements of $P(D)$ must be between 0 and 1 and sum to 1


- Solutions

- Use subsets of $\mathbf{S}$; average results
- Equivalent to kernel density smoothing of sparse categorical data
- Use constrained LS to constrain $P(D)$ to simplex


- Uncertainty estimates by bootstrapping

A Nonrandom Hand-coded Sample

**Differences in Document
Category Frequencies**




**Differences in Word
Profile Frequencies**




All existing methods would fail with these data.

Accurate Estimates




Out of Sample Validation: Blogs




Out of Sample Validation: Other Examples

Movie Reviews




University Websites




Bias by Number of Hand Coded Documents


Nonparametric Estimator



Sampling Estimator




Average RMSE by Number of Hand Coded Documents




Misclassification Matrix for Blog Posts

	-2	-1	0	1	2	NA	NB	$P(D_1)$
-2	.70	.10	.01	.01	.00	.02	.16	.28
-1	.33	.25	.04	.02	.01	.01	.35	.08
0	.13	.17	.13	.11	.05	.02	.40	.02
1	.07	.06	.08	.20	.25	.01	.34	.03
2	.03	.03	.03	.22	.43	.01	.25	.03
NA	.04	.01	.00	.00	.00	.81	.14	.12
NB	.10	.07	.02	.02	.02	.04	.75	.45


SIMEX Analysis of “Not a Blog” Category









SIMEX Analysis of “Not a Blog” Category



SIMEX Analysis of “Not a Blog” Category



SIMEX Analysis of Other Categories



What can go wrong?

What can go wrong?

- We assume $P^h(\mathbf{S}|D) = P(\mathbf{S}|D)$

What can go wrong?

- We assume $P^h(\mathbf{S}|D) = P(\mathbf{S}|D)$
- Must choose word stem subset size (a smoothing parameter)

What can go wrong?

- We assume $P^h(\mathbf{S}|D) = P(\mathbf{S}|D)$
- Must choose word stem subset size (a smoothing parameter)
- Need enough labeled documents in each category (can hand code more if CI's are too large, perhaps via case-control methods)

What can go wrong?

- We assume $P^h(\mathbf{S}|D) = P(\mathbf{S}|D)$
- Must choose word stem subset size (a smoothing parameter)
- Need enough labeled documents in each category (can hand code more if CI's are too large, perhaps via case-control methods)
- Need sufficient information in: documents, categorization scheme, numerical summaries of the documents, and hand-codings

What can go wrong?

- We assume $P^h(\mathbf{S}|D) = P(\mathbf{S}|D)$
- Must choose word stem subset size (a smoothing parameter)
- Need enough labeled documents in each category (can hand code more if CI's are too large, perhaps via case-control methods)
- Need sufficient information in: documents, categorization scheme, numerical summaries of the documents, and hand-codings
- Use additional hand coding to verify assumptions

Verbal Autopsy Methods

Verbal Autopsy Methods

- The Problem

Verbal Autopsy Methods

- The Problem
 - Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies

- The Problem

- Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies
- High quality death registration: only 23/192 countries

Verbal Autopsy Methods

- The Problem
 - Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies
 - High quality death registration: only 23/192 countries
- Existing Approaches

Verbal Autopsy Methods

- The Problem
 - Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies
 - High quality death registration: only 23/192 countries
- Existing Approaches
 - Ask relatives or caregivers 50-100 symptom questions

- The Problem

- Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies
- High quality death registration: only 23/192 countries

- Existing Approaches

- Ask relatives or caregivers 50-100 symptom questions
- Ask physicians to determine cause of death (low intercoder reliability)

- The Problem

- Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies
- High quality death registration: only 23/192 countries

- Existing Approaches

- Ask relatives or caregivers 50-100 symptom questions
- Ask physicians to determine cause of death (low intercoder reliability)
- Apply expert algorithms (high reliability, low validity)

- The Problem

- Policymakers need the **cause-specific mortality rate** to set research goals, budgetary priorities, and ameliorative policies
- High quality death registration: only 23/192 countries

- Existing Approaches

- Ask relatives or caregivers 50-100 symptom questions
- Ask physicians to determine cause of death (low intercoder reliability)
- Apply expert algorithms (high reliability, low validity)
- Find deaths with medically certified causes from a local hospital, trace caregivers to their homes, ask the same symptom questions, and statistically classify deaths in population (model-dependent, low accuracy)

An Alternative Approach

An Alternative Approach

- Document-Category, Cause of Death,

$$D_i = \begin{cases} 1 & \text{if bladder cancer} \\ 2 & \text{if cardiovascular disease} \\ 3 & \text{if transportation accident} \\ \vdots & \vdots \\ J & \text{if infectious respiratory} \end{cases}$$

An Alternative Approach

- Document ~~Category~~, Cause of ~~Death~~,

$$D_i = \begin{cases} 1 & \text{if bladder cancer} \\ 2 & \text{if cardiovascular disease} \\ 3 & \text{if transportation accident} \\ \vdots & \vdots \\ J & \text{if infectious respiratory} \end{cases}$$

- Word ~~Stem~~ Profile, ~~Symptoms~~:

$$S_i = \begin{cases} S_{i1} = 1 & \text{if "breathing difficulties", 0 if not} \\ S_{i2} = 1 & \text{if "stomach ache", 0 if not} \\ \vdots & \vdots \\ S_{iK} = 1 & \text{if "diarrhea", 0 if not} \end{cases}$$

An Alternative Approach

- Document ~~Category~~, Cause of ~~Death~~,




$$D_i = \begin{cases} 1 & \text{if bladder cancer} \\ 2 & \text{if cardiovascular disease} \\ 3 & \text{if transportation accident} \\ \vdots & \vdots \\ J & \text{if infectious respiratory} \end{cases}$$

- Word ~~Stem~~ Profile, ~~Symptoms~~:



$$S_i = \begin{cases} S_{i1} = 1 & \text{if "breathing difficulties", 0 if not} \\ S_{i2} = 1 & \text{if "stomach ache", 0 if not} \\ \vdots & \vdots \\ S_{iK} = 1 & \text{if "diarrhea", 0 if not} \end{cases}$$

- Apply the **same** methods

Validation in China



Validation in Tanzania



For more information

<http://GKing.Harvard.edu>